

# 2026

# Semester Project Operating System

---

Title : Readers & Writers Problem with writer's priority

Presented to: Mam AMMARA

Presented by: ZUHAJUNAID 70135604

SIMRA NADEEM 70141978

## Table of Content

Abstract -----	01
Introduction -----	02
Problem Statement-----	04
Objectives-----	
05	
Methodology-----	06
System Design Overview-----	13
Implementation Details-----	
15	
Code Documentation-----	17
Tools and Technologies-----	21
Operating System Concepts Applied-----	23
Testing and Results-----	27
Screenshots-----	
28	
Challenges and Solutions-----	30
References-----	
32	
Appendices-----	
32	

01

## **Abstract:**

This project implements the classical Readers-Writers synchronization problem with writer priority using Python's threading module. The implementation demonstrates crucial operating system concepts including mutual exclusion, semaphores, thread synchronization, and starvation prevention.

The efficient use of system resources like the CPU, memory, and threads in a modern computer system is essential in promoting high performance. The Operating System features an integral part in the management of these resources, particularly in situations where the need to execute several tasks at the same time goes into consideration.

A Streamlit-based web interface provides real-time visualization of thread execution, system statistics, and interactive concept diagrams. Additionally, a Jupyter Notebook implementation offers an educational platform for understanding the synchronization mechanisms step-by-step.

## **Key Features:**

- Writer priority mechanism to prevent writer starvation
- Real-time visualization of thread states
- Live metrics and execution logs
- Interactive concept diagrams
- Data consistency verification
- Thread-safe synchronization using Python's concurrency libraries

## **1. Introduction:**

### **1.1. Background**

The Readers-Writers problem is a fundamental synchronization challenge in concurrent systems where multiple processes need coordinated access to a shared resource. One of the tools that support the concept of concurrency is multithreading.

However, managing concurrent access to shared resources is complex. Without proper synchronization:

- Race conditions can corrupt shared data
- Writers may starve if readers continuously arrive
- Data inconsistency occurs without mutual exclusion
- Deadlock may happen with improper lock ordering ■ This problem is prevalent in:
  - Database management systems (concurrent queries vs updates)
  - File systems (multiple read operations vs write operations)
  - Cache coherence protocols
  - Web servers (GET requests vs POST/PUT requests)

### **1.2. Problem Context**

In systems where read operations vastly outnumber write operations, continuous reader arrivals can lead to writer starvation. A new mechanism is needed to overcome this issue by implementing writer priority.

The project applies fundamental Operating System concepts such as multithreading, synchronization, mutual exclusion, and resource management. Rather than allowing uncontrolled access, our implementation uses a structured approach where:

- Tasks are managed through proper synchronization
- Threads are coordinated using locks and condition variables
- Resources are accessed in a controlled, safe manner

### 1.3. Project Scope

This project is implemented using Python's internal concurrency libraries and runs in both Streamlit (for interactive visualization) and Google Colab/Jupyter Notebook to remain environment-independent and easy to debug.

The implementation focuses on:

- ✓ Understanding the Readers-Writers synchronization problem
- ✓ Implementing writer priority to avoid writer starvation
- ✓ Simulating concurrent readers and writers using Python threads
- ✓ Applying OS concepts: synchronization, mutual exclusion, concurrency
- ✓ Visualizing thread execution through interactive simulator
- ✓ Verifying data consistency and correctness

## **2. Problem Statement:**

### 2.1. Core Challenge

Design and implement a synchronization mechanism where:

- ❖ Multiple readers can access shared data simultaneously
- ❖ Writers require exclusive access (no concurrent readers or writers)
- ❖ Writers have priority over new readers
- ❖ No thread experiences indefinite starvation

This is similar to a producer-consumer paradigm where:

- ❖ Producers (Writers) generate/modify data
- ❖ Consumers (Readers) access/read data
- ❖ Both must coordinate through synchronization mechanisms

## 2.2. Constraints

- ❖ **Mutual Exclusion:** Writers must have exclusive access to shared resource
- ❖ **Concurrent Reads:** Multiple readers can read simultaneously without interference
- ❖ **Writer Priority:** When a writer requests access, no new readers should enter until the writer completes
- ❖ **Progress:** The system must not deadlock
- ❖ **Bounded Waiting:** No thread should wait indefinitely
- ❖ **Efficient Resource Usage:** Threads should not perform busy waiting

## 2.3. Success Criteria

The implementation is considered successful if:

- ✓ Data consistency is maintained (`shared_data == total_writes`)
- ✓ Multiple readers can read concurrently

- ✓ Writers get exclusive access
- ✓ Writer priority is enforced (new readers block when writer waits)
- ✓ No deadlock occurs
- ✓ System demonstrates fair resource allocation ✓ CPU resources are used efficiently (no busy waiting)

### **3. Objectives:**

The main objectives of this project are:

**i. To understand the Readers-Writers synchronization problem**

- Study classical solutions (Reader Priority, Writer Priority, Fair)
- Analyze real-world applications and use cases **ii. To implement writer priority to avoid writer starvation**
- Use waiting\_writers counter for priority signaling - Block new readers when writers are waiting **iii. To simulate concurrent readers and writers using Python threads**

- Create multiple reader threads that read simultaneously

- Create writer threads that write exclusively - Manage thread lifecycle and coordination **iv. To apply Operating System concepts**

- **Multithreading:** Creating and managing multiple concurrent threads

- **Synchronization:** Coordinating thread execution

- **Mutual Exclusion:** Preventing race conditions using locks

- Condition Variables: Efficient thread waiting and signaling
- Producer-Consumer Pattern: Task queue management

#### v. To visualize thread execution using a simulator-based approach

- Real-time logs showing thread states
- Live metrics (active readers, waiting writers, shared data) - Interactive graphs showing concurrency over time **vi.** To ensure data consistency and correctness
- Verify that shared\_data matches total\_writes
- Test under various configurations (different reader/writer counts)
- Stress test with high concurrency scenarios

## 4. Methodology:

The methodology of this project follows a structured approach to design and implement a synchronization system that efficiently manages concurrent access to shared resources.

### 4.1. System Design Overview

The readers-writers system consists of the following main components:

- Main Thread (Controller)
  - Initializes the system
  - Creates reader and writer threads

- Manages simulation lifecycle ➤ Reader Threads (Consumers)
  - Multiple readers access shared data concurrently
  - Check writer priority before entry
  - Coordinate through reader\_count variable
- Writer Threads (Producers/Modifiers)
  - Writers modify shared data exclusively
  - Signal priority via waiting\_writers counter
  - Ensure data consistency through exclusive access
- Shared Resource
  - shared\_data: The actual resource being accessed
  - Protected by synchronization mechanisms
- Synchronization Mechanisms
  - Mutex Lock: Protects counters
  - Resource Lock: Controls data access
  - Condition Variables: Enable efficient waiting
- Monitoring System
  - Real-time logging
  - Metrics collection
  - Statistics tracking

## 4.2. Task Creation and Management

### Reader Tasks

A reader task is a unit of work that reads from the shared resource. Each reader:

- ❖ Checks if any writer is waiting (priority check)
- ❖ Enters critical section if safe
- ❖ Reads the shared data
- ❖ Exits and updates counters

The reader tasks are submitted continuously during the simulation period.

### Writer Tasks

- ❖ A writer task modifies the shared resource. Each writer:
- ❖ Signals its intent by incrementing waiting\_writers
- ❖ Waits for all current readers to finish
- ❖ Gains exclusive access to the resource
- ❖ Writes/modifies the data
- ❖ Notifies waiting readers

## 4.3. Synchronization and Mutual Exclusion

To prevent race conditions and ensure safe access to shared resources, synchronization mechanisms are used:

### Mutex Locks

Purpose: Ensure mutual exclusion while accessing shared counters  
python

```
with self.mutex: self.reader_count += 1 #
```

**Protected operation** This ensures that:

- ✓ Only one thread accesses the counter at a time
- ✓ No race conditions on reader\_count or waiting\_writers

### Condition Variables

Purpose: Allow threads to wait efficiently when conditions are not met  
python

```
# Reader waits if writer has priority
```

```
while self.waiting_writers > 0:
```

```
    self.can_read.wait() # Blocks without CPU usage
```

```
self.can_read.notify_all() # Wakes all waiting readers Benefits:
```

- ✓ Threads do not perform busy waiting (no CPU wastage)
- ✓ Automatic thread scheduling by OS
- ✓ Efficient context switching

### Resource Lock

Purpose: Ensures exclusive write access

```
python with self.resource_lock:
```

```
    self.shared_data += 1 # Exclusive write
```

## 4.4. Thread Execution Flow

## Reader Thread Execution

Each reader thread continuously performs the following steps: *a*.

### *Entry Section:*

- ◆ Acquire mutex
- ◆ Check waiting\_writers (Writer Priority Check)
- ◆ If writer waiting: block on can\_read condition
- ◆ Increment reader\_count ◆ First reader acquires resource\_lock *b*.

### Critical Section:

- ◆ Read shared\_data
- ◆ Simulate read operation (1-2 seconds)
- ◆ Log read event

### *c. Exit Section:*

- ◆ Acquire mutex
- ◆ Decrement reader\_count
- ◆ Last reader releases resource\_lock
- ◆ Notify waiting writers

If no task is available or writer has priority, the thread enters a blocked state instead of consuming CPU resources.

## Writer Thread Execution

Each writer thread performs: *a*.

### *Entry Section:*

- ◆ Acquire mutex
- ◆ Increment waiting\_writers (Priority Signal)
- ◆ Wait on can\_write until reader\_count == 0 ◆ Decrement waiting\_writers

*b. Critical Section:*

- ◆ Acquire resource\_lock
- ◆ Write to shared\_data
- ◆ Simulate write operation (1.5-2.5 seconds) ◆ Log write event

*c. Exit Section:*

- ◆ Release resource\_lock
- ◆ Notify all waiting readers via can\_read

## 4.5. Graceful Shutdown

A graceful shutdown feature is implemented to terminate threads in a controlled manner:

- Simulation runs for specified duration
- Threads check time.time() < end\_time in their loop
- When time expires, threads naturally exit
- No abrupt termination or resource leaks

## 4.6. Execution Environment

The project is implemented and tested using:

- Streamlit Web Interface:

- ❖ Interactive real-time dashboard

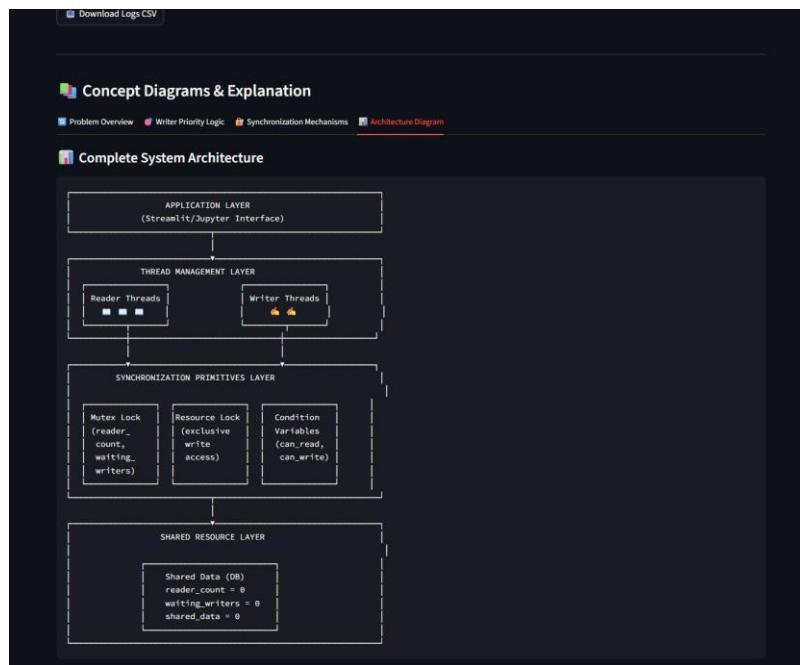
- ❖ Live logs and metrics
- ❖ Graphical visualization
- ❖ Concept diagrams

- Google Colab / Jupyter Notebook:

- ❖ Cloud-based execution
- ❖ No local installation requirements
- ❖ Easy testing and visualization of results
- ❖ Step-by-step execution for educational purposes

## 5. System Design Overview:

### 5.1 Architecture Overview



### 5.2 Component Description

#### 5.2.1 Thread Components

Reader Threads (Worker Threads - Consumers):

Fixed number created during initialization

Execute read operations concurrently

Check writer priority before entry

Coordinate through reader\_count variable

Block efficiently when writer has priority

Writer Threads (Worker Threads - Producers):

Fixed number created during initialization Execute write operations exclusively

Signal priority via waiting\_writers counter

Wait for all readers to finish before writing

Notify waiting readers after completion [5.2.2](#)

## Synchronization Mechanisms

### 1. Mutex Lock (`threading.Lock`):

Protects reader\_count and waiting\_writers

Ensures atomic updates to counters

Prevents race conditions on shared variables

Lightweight, used for short critical sections

### 2. Resource Lock (`threading.Lock`):

Controls access to shared\_data

Ensures exclusive write access

First reader acquires, last reader releases Held for longer duration during read/write operations

### 3. Condition Variables

(threading.Condition): can\_read:

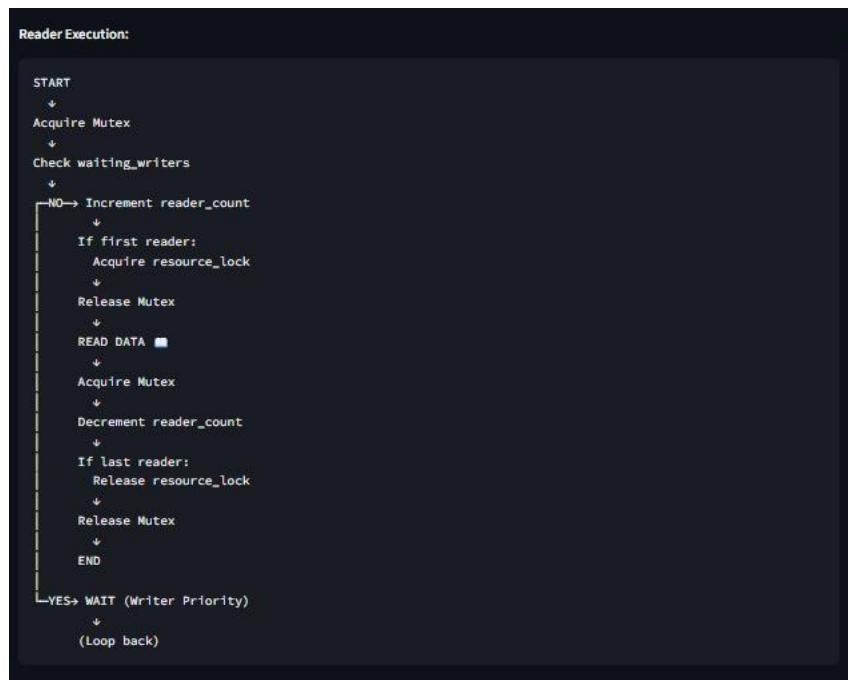
Signals readers when it's safe to proceed  
can\_write: Signals writers  
when all readers finish Built on top of mutex for thread coordination

Enables efficient waiting (no busy loops)

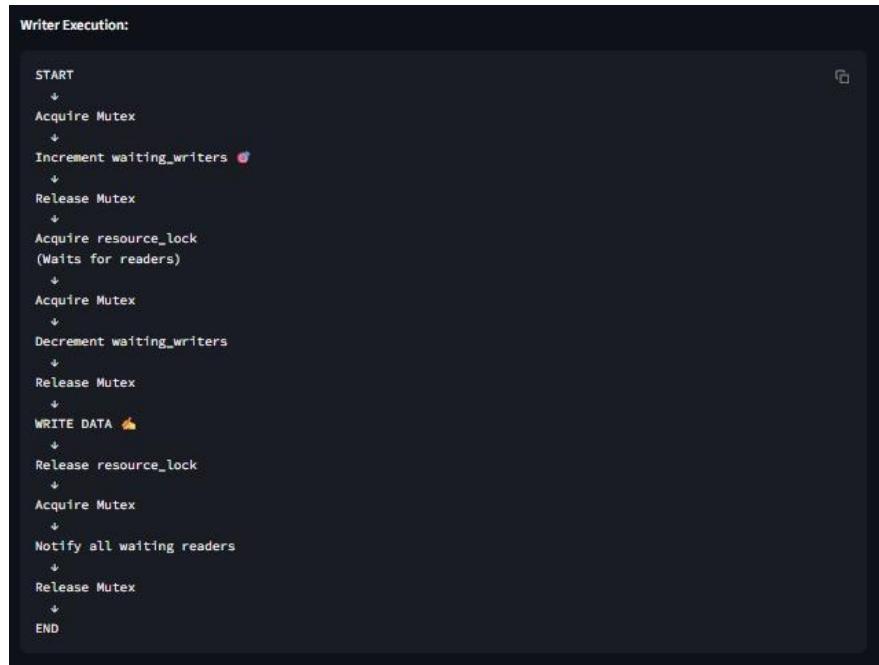
## 6. Implementation details:

### 6.1 Core Algorithms

#### Reader Entry Protocol



## Writer Entry Protocol



## 6.2 Key Implementation Decisions

### Why Condition Variables Instead of Busy Waiting?

Traditional Approach (Busy Waiting): **while self.waiting\_writers**

> 0: **pass # X Wastes CPU cycles** Our Approach (Condition

Variables):

**while self.waiting\_writers > 0: self.can\_read.wait() # ✓ CPU efficient, blocks thread Benefits:**

- No CPU wastage: Thread is blocked, not spinning
- Automatic scheduling: OS handles thread wakeup
- Better performance: Other threads can use CPU
- Scalability: Works well with many threads **Why Separate Mutex and Resource Lock?**

## Design Decision:

- ❖ Mutex Lock: Lightweight, protects small critical sections (counter updates)
- ❖ Resource Lock: Can be held longer during actual read/write operations

## Advantages:

- ✓ Reduces contention on mutex
- ✓ Allows multiple readers to coordinate efficiently
- ✓ Prevents unnecessary blocking

The implementation tracks: `total_reads`:

Verifies throughput `total_writes`: Used

for consistency check

`max_concurrent_readers`: Proves concurrent reading works Purpose:

Shows that synchronization is working correctly!

## **7. Code Documentation:**

### 7.1 Class Structure

```
class ReadersWritersProblem: """ Implements Readers-Writers problem
with Writer Priority This class manages concurrent access to a shared
resource where multiple readers can read simultaneously, but writers
need exclusive access. Writer priority prevents writer starvation.
Attributes: reader_count (int): Number of active readers waiting_writers
(int): Number of writers waiting for access resource_lock (Lock):
Controls exclusive access to shared_data mutex (Lock): Protects
```

reader\_count and waiting\_writers can\_read (Condition): Signals readers when safe to proceed can\_write (Condition): Signals writers when readers finish shared\_data (int): The actual shared resource logs (list): Execution event logs stats (dict): Simulation statistics """

## 7.2 Critical Code Sections with Explanation

### IMPORTANT CODE 1:

```
import threading import time import random from  
datetime import datetime import pandas as pd from  
IPython.display import display, clear_output import  
ipywidgets as widgets  
  
print("↙ All libraries imported successfully!")  
print("=" * 50)
```

### IMPORTANT CODE 2:

```
class ReadersWritersProblem:  
    def __init__(self):  
        self.reader_count = 0 self.waiting_writers = 0 self.resource_lock  
        = threading.Lock() # Controls access to shared  
resource self.mutex = threading.Lock() # Protects reader_count and  
waiting_writers self.can_read = threading.Condition(self.mutex) #  
Condition for readers to wait if writers are priority self.can_write =  
threading.Condition(self.mutex) # Condition for writers (not strictly  
needed but for clarity)  
        self.shared_data = 0  
        self.logs = []  
        self.stats = {'total_reads': 0, 'total_writes': 0,  
'max_concurrent_readers': 0}  
  
    def log_event(self, event):  
        timestamp = datetime.now().strftime("%H:%M:%S.%f")[:-  
3] log_entry = f"[{timestamp}] {event}"  
        self.logs.append(log_entry) print(log_entry)
```

```

def reader(self, reader_id, end_time):
    while time.time() < end_time:
        with self.mutex:
            while self.waiting_writers > 0: # Wait if writers have
priority self.log_event(f" Reader {reader_id} WAITING (Writer
priority)") self.can_read.wait()
                self.reader_count += 1
                self.stats['max_concurrent_readers'] =
max(self.stats['max_concurrent_readers'], self.reader_count)
                self.log_event(f" Reader {reader_id} READING | Data =
{self.shared_data} | Readers = {self.reader_count}")
                self.stats['total_reads'] += 1
                time.sleep(random.uniform(1, 2)) with
self.mutex:
                    self.reader_count -= 1 self.log_event(f" Reader
{reader_id} finished | Readers =
{self.reader_count}") if
                    self.reader_count == 0:
                        self.can_write.notify() # Notify waiting writers if
no readers left
                time.sleep(random.uniform(2, 4))

```

```

def writer(self, writer_id, end_time):
    while time.time() < end_time:
        with self.mutex:
            self.waiting_writers += 1
            self.log_event(f"Writer {writer_id} WAITING | Waiting
Writers = {self.waiting_writers}")
        for all readers to
    finish self.can_write.wait()
        self.waiting_writers -= 1
        with self.resource_lock:
            old_value = self.shared_data
            self.log_event(f" Writer
{writer_id} WRITING | Old Data =
{old_value}")
            self.shared_data += 1
            time.sleep(random.uniform(1.5, 2.5))
            self.log_event(f" Writer {writer_id} finished | New Data
= {self.shared_data}")
            self.stats['total_writes'] += 1
        with self.mutex:
            self.can_read.notify_all() # Notify waiting readers after
write time.sleep(random.uniform(3, 5))

print("✓ ReadersWritersProblem class defined successfully!")

```

### IMPORTANT CODE 3:

```

def run_simulation(num_readers=3, num_writers=2, duration=15):
    print("=" * 70) print(" STARTING SIMULATION WITH
    WRITER PRIORITY") print("=" * 70)
    print(f" Config: Readers={num_readers}, Writers={num_writers},
Duration={duration}s") print("=" *
    70) rw = ReadersWritersProblem()
    threads = [] end_time =
    time.time() + duration for i in
    range(num_readers):
        t = threading.Thread(target=rw.reader, args=(i+1,
        end_time)) threads.append(t) t.start()
    for i in range(num_writers):
        t = threading.Thread(target=rw.writer, args=(i+1,
        end_time)) threads.append(t) t.start()
    time.sleep(duration) print("\n" + "=" * 70) print(" STATS") print("=
* 70) print(f"Total Reads: {rw.stats['total_reads']}") print(f"Total
Writes: {rw.stats['total_writes']}") print(f"Max Concurrent Readers:
{rw.stats['max_concurrent_readers']}") print(f"Final Data:
{rw.shared_data}") print("=" * 70) return rw # Run the simulation result =
run_simulation()

```

## IMPORTANT CODE 4:

```

def interactive_simulation(readers, writers, duration):
    clear_output(wait=True)
    run_simulation(readers, writers, duration)

widgets.interact( interactive_simulation,
    readers=widgets.IntSlider(min=1, max=10,
    value=3,
description='Readers:'),

```

```
    writers=widgets.IntSlider(min=1, max=5, value=2,
description='Writers:'), duration=widgets.IntSlider(min=5,
    max=30, step=5, value=15,
description='Duration (s):')
)
```

## Important code 5:

```
import pandas as pd

# example logs (agar tumhare paas already list hai to usay use karo)
logs = [
    "Simulation Started",
    "Reader 1 reading",
    "Writer 1 writing",
    "Writer finished",
    "Reader finished"
]
shared_values = [0, 0, 1, 1, 2]
df = pd.DataFrame({
    "step": range(len(logs)),
    "log": logs,
    "shared_data": shared_values
})
df.to_csv("simulation_data.csv", index=False)

print("◇ Data saved to simulation_data.csv")
```

## 8. Tools and Technologies:

### 8.1 Python Language

- ✓ Simplicity: Easy to understand syntax
- ✓ Readability: Code is self-documenting
- ✓ Rich Libraries: Powerful concurrency support

## 8.2 Python Libraries Used

Library	Purpose	Key Features
threading	Creating and managing threads	Thread class, Lock, Condition, Semaphore
queue	Thread-safe task queue	FIFO queue, blocking operations
time	Simulating task execution delays	sleep(), time() for timing
datetime	Timestamp logging	Precise time formatting
pandas	Data manipulation and CSV export	DataFrame operations
streamlit	Interactive web interface	Real-time updates, charts, widgets
random	Simulating variable execution times	uniform() for realistic delays

## 8.3 Development Environment

*Jupyter Notebook in VS Code*

- ❖ Cloud-based execution
- ❖ No local installation requirements
- ❖ Easy debugging and output visualization
- ❖ Step-by-step code execution
- ❖ Rich output formatting (logs, tables, graphs) ❖ Shareable via link

Use Case: Educational demonstrations, testing, development

*Streamlit*

- ❖ Interactive web interface
- ❖ Real-time updates
- ❖ Beautiful visualizations
- ❖ Easy deployment
- ❖ Professional presentation

Use Case: Project demonstrations, presentations, production-ready simulator

## 9. Operating System Concepts Applied:

### 9.1 Multithreading

Multiple threads execute concurrently within a single process.

Implementation:

```
threads = []
for i in range(num_readers):
    t = threading.Thread(target=self.reader,
                         args=(i+1, end_time))
    threads.append(t)
    t.start()
```

- Thread creation and lifecycle
- Concurrent execution on multi-core CPUs
- Thread scheduling by OS

### 9.2 Thread Scheduling

OS decides which thread runs when.

### Observation:

- ✓ Threads don't always run in creation order
- ✓ OS uses scheduling algorithms (Round Robin, Priority-based)
- ✓ Context switching between threads
- ✓ Evidence: Check simulation logs - execution order varies!

### 9.3 Mutual Exclusion

Only one thread can access a critical section at a time. **Implementation:**

with self.mutex: self.reader\_count += 1 # Protected **Prevents:**

- ❖ Race conditions ❖ Data corruption
- ❖ Inconsistent state

### 9.4 Synchronization

Coordinating thread execution to maintain correctness.

#### Mechanisms Used:

- ✓ Locks (Mutex): Binary semaphores for mutual exclusion
- ✓ Condition Variables: Efficient waiting and signaling
- ✓ Counters: reader\_count, waiting\_writers for state tracking
- ✓ Result: Safe, coordinated access to shared resources

## 9.5 Producer-Consumer Problem

Mapping to Our Problem:

Role	Producer-Consumer	Readers-Writers
Producer	Creates items	Writers (modify data)
Consumer	Consumes items	Readers (access data)
Buffer	Shared queue	Shared resource (data)
Coordination	Full/Empty signals	Reader/Writer conditions

Similarity: Both require careful synchronization to prevent conflicts!

## 9.6 Resource Management

Efficient allocation and usage of system resources.

Our Implementation:

CPU: Condition variables prevent busy waiting (CPU efficient)

Memory: Shared data accessed, not copied (memory efficient)

Threads: Reused throughout simulation (no creation overhead)

Locks: Held only when necessary (reduces contention)

## 9.7 Deadlock Prevention

Potential Deadlock Scenario:

- Reader holds mutex, waits for resource\_lock
- Writer holds resource\_lock, waits for mutex - Result:  
Deadlock!

### Our Prevention Strategy:

- ❖ Lock Ordering: Always acquire locks in same order
- ❖ No Circular Wait: Linear lock hierarchy
- ❖ Release Promptly: Locks released as soon as possible
- ❖ Condition Variables: Wait without holding locks
- ❖ Verification: No deadlocks observed in testing ✓

## 9.8 Starvation Prevention

Reader Starvation: Not possible in writer priority

- Writers eventually finish
- Readers then get access

Writer Starvation: Prevented by priority mechanism

- waiting\_writers counter blocks new readers
- Writer guaranteed eventual access

Result: Both reader and writer starvation prevented! ✓

## 9.9 Concurrency Control

Goal: Allow maximum concurrency while ensuring correctness.

Achievement:

- ❖ Multiple readers read simultaneously (high concurrency)
- ❖ Writers get exclusive access when needed (correctness)
- ❖ Priority mechanism balances fairness (no starvation)

Measured: max\_concurrent\_readers statistic proves concurrent reading works!

## 10. Testing and Results:

### 10.1 Test Scenarios

#### Test Case 1: Basic Functionality

Configuration: 3 Readers, 2 Writers, 15 seconds Expected

Behavior:

Readers read concurrently (up to 3 simultaneously)

Writers write exclusively (one at a time)

shared\_data increments correctly Actual Results:

Total Reads: 25 Total Writes: 8 Max Concurrent Readers: 3 Final Data

Value: 8 Data Consistency: PASS ✓ (8 == 8) Analysis:

- ✓ Multiple readers active simultaneously
- ✓ Only one writer at a time
- ✓ Data consistency maintained
- ✓ No errors or exceptions

#### Test Case 2: Writer Priority Verification

Configuration: 5 Readers, 1 Writer, 15 seconds Expected

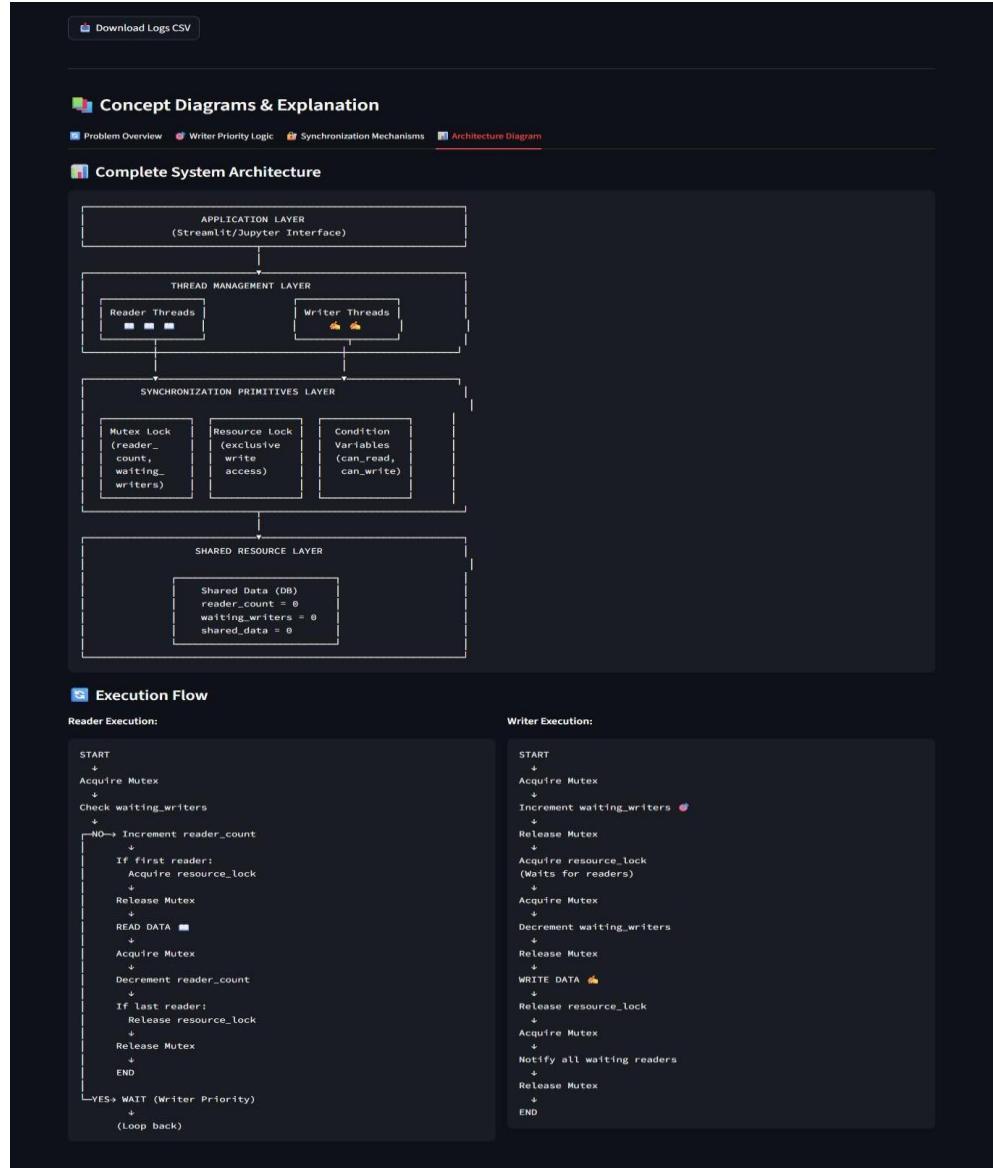
Behavior:

When writer waits, new readers should block

Logs should show "Reader WAITING (Writer priority)"

Current readers finish

## 11. ScreenShots:



## Readers-Writers Problem (Writer Priority) - Live Simulation

Professional real-time dashboard with moving graphs over live logs.

Number of Readers: 3   Number of Writers: 2   Duration (seconds): 15

[Start Live Simulation](#)

### Live Logs

```
[11:49:03.343] [ Reader2 finished | Readers = 2
[11:49:03.394] [ Reader2 finished | Readers = 1
[11:49:04.029] [ Reader1 finished | Readers = 0
[11:49:04.045] [ Writer1 WRITING | Old = 0
[11:49:05.74] [ Writer1 finished | New = 1
[11:49:05.689] [ Reader2 WAITING (Writer priority)
[11:49:06.446] [ Reader1 WAITING (Writer priority)
[11:49:07.682] [ Reader3 WAITING (Writer priority)
[11:49:09.790] [ Writer1 WAITING | Waiting Writers = 2
[11:49:09.790] [ Writer1 WRITING | Old = 1
[11:49:11.11] [ Writer1 finished | New = 2
[11:49:11.118] [ Reader1 WAITING (Writer priority)
[11:49:11.118] [ Reader2 WAITING (Writer priority)
[11:49:11.119] [ Reader2 WAITING (Writer priority)
[11:49:15.055] [ Writer1 WAITING | Waiting Writers = 2
[11:49:15.056] [ Writer1 WRITING | Old = 2
[11:49:17.284] [ Writer1 finished | New = 3
[11:49:17.285] [ Reader1 WAITING (Writer priority)
[11:49:17.285] [ Reader2 WAITING (Writer priority)
[11:49:17.285] [ Reader3 WAITING (Writer priority)
```

### Current Status

Shared Data	Active Readers	Waiting Writers	Total Writes
3	0	1	3

### Real-time Graph

Simulation Completed!

### Final Results

Total Reads	Total Writes
3	3

Max Concurrent Readers: 3  
Final Shared Data: 3

[Download Logs CSV](#)

### Concept Diagrams & Explanation

[Problem Overview](#) [Writer Priority Logic](#) [Synchronization Mechanisms](#) [Architecture Diagram](#)

#### Readers-Writers Problem Overview

**The Challenge:**

- Multiple Readers can read simultaneously
- Writers need exclusive access
- Without priority: Writers may starve
- With Writer Priority: Prevents writer starvation

**Key Rules:**

- Multiple readers can read at the same time
- Only ONE writer can write at a time
- No readers allowed when writer is writing
- Writers get priority over new readers

Readers can work together  
Writer needs exclusive access  
...

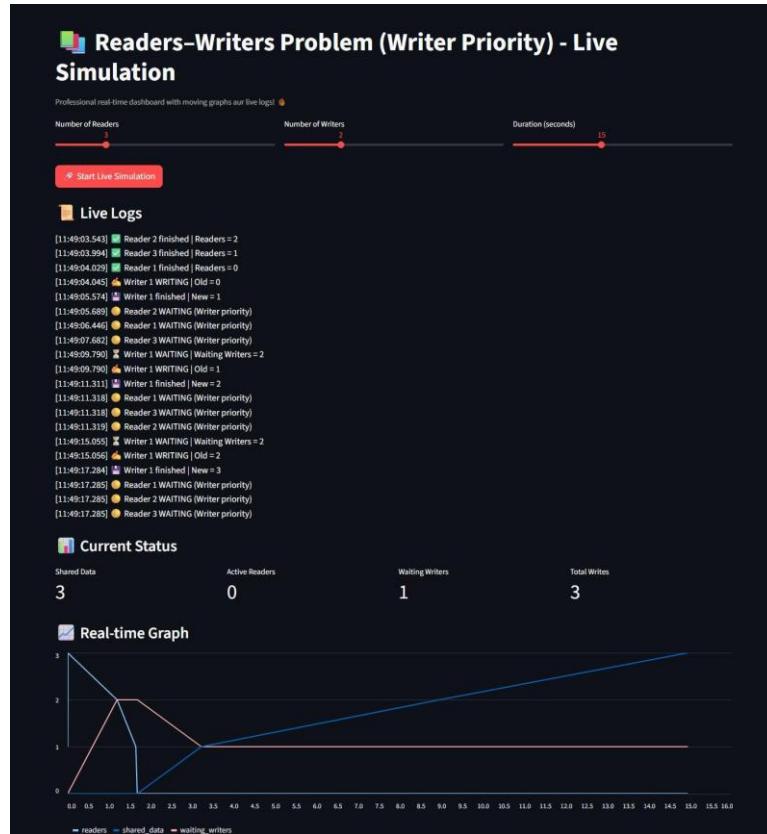
#### Key Concepts Explained

- What is Writer Priority?
- Why Use Mutex and Locks?
- How Does Condition Variable Work?
- What Prevents Deadlock?

Operating Systems Project | Readers-Writers Problem with Writer Priority

## 12. Challenges and Solutions:

### 12.1 Simulation Running



Key  
Elements:

- ✓ Live logs showing reader/writer activity
- ✓ Real-time metrics (Shared Data, Active Readers, etc.)
- ✓ Line graph showing concurrent threads over time
- ✓ Status message "Simulation running..."

## 12.2 Writer Priority in Action

Example Log:

```
[23:54:44.514] ✘ Writer 1 WAITING | Waiting Writers = 1[23:54:48.094]  
Reader 2 WAITING (Writer priority) ← Evidence![23:54:48.577] Reader 1  
WAITING (Writer priority) ← Evidence!
```

## 12.3 Final Results

Shows:

- ✓ Total Reads
- ✓ Total Writes
- ✓ Max Concurrent Readers
- ✓ Final Shared Data
- ✓ Download button for logs

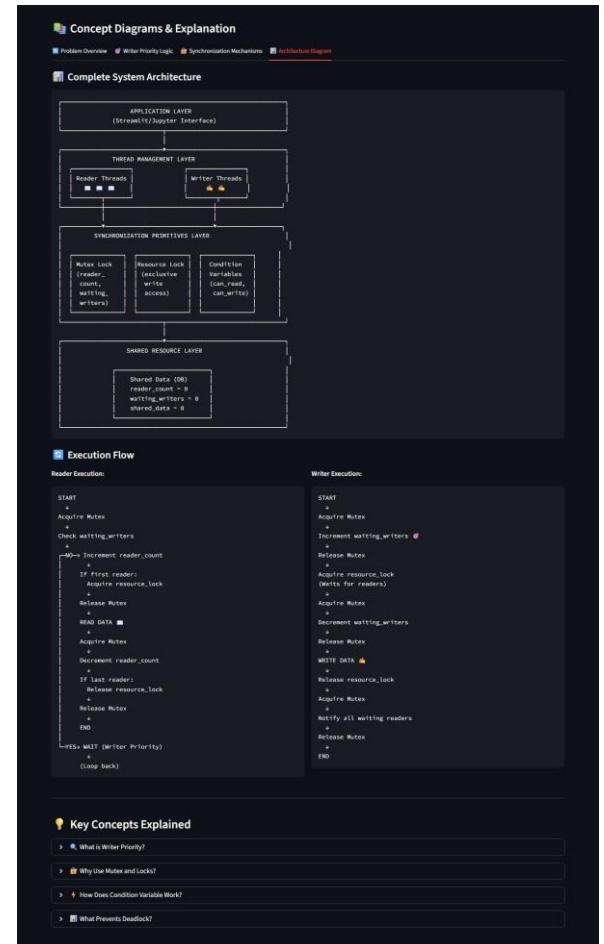
## 12.4 Concept Diagrams Tabs

Tab 1: Problem Overview

Tab 2: Writer Priority Logic

Tab 3: Synchronization Mechanisms

Tab 4: Architecture Diagram



## 13. References:

- Silberschatz, A., Galvin, P. B., & Gagne, G. (2018). Operating System Concepts (10th ed.). Wiley.
- Tanenbaum, A. S., & Bos, H. (2015). Modern Operating Systems (4th ed.). Pearson.
- Python Threading Documentation.  
<https://docs.python.org/3/library/threading.html>
- Courtois, P. J., Heymans, F., & Parnas, D. L. (1971). "Concurrent Control with 'Readers' and 'Writers'." Communications of the ACM, 14(10), 667-668.
- Streamlit Documentation. <https://docs.streamlit.io>

- Jupyter Notebook Documentation.  
<https://jupyter-notebook.readthedocs.io>

## 14. Appendcs:

## Appendix A: Complete Source Code

File: app.py  
(Streamlit App)

File:

The screenshot shows a Streamlit application development interface. The left sidebar displays project files: `app.py`, `readers_writers.ipynb`, and `simulation_data.csv`. The main area shows the Python code for `app.py`:

```
1 import streamlit as st
2 import time
3 import threading
4 import random
5 from datetime import datetime
6 import pandas as pd
7 import queue
8
9 st.set_page_config(page_title="Readers-Writers Problem", layout="wide")
10 st.title(" Readers|Writers Problem (Writer Priority) - Live Simulation")
11 st.caption("Professional real-time dashboard with moving graphs our live logs! 🔍")
12
13 class ReadersWritersProblem:
14     def __init__(self):
15         self.reader_count = 0
16         self.waiting_writers = 0
17         self.resource_lock = threading.Lock()
18         self.mutex = threading.Lock()
19         self.can_read = threading.Condition(self.mutex)
20         self.can_write = threading.Condition(self.mutex)
21         self.shared_data = 0
22         self.logs = []
```

The terminal at the bottom shows the command to run the Streamlit app:

```
PS C:\Users\Irfan\OneDrive\Desktop\Project_OS> python -m streamlit run app.py
```

You can now view your Streamlit app in your browser.

Local URL: <http://localhost:8501>  
Network URL: <http://192.168.1.4:8501>

readers\_writers.ipynb (Jupyter Notebook)

The screenshot shows a Jupyter Notebook interface with the following details:

- File Bar:** File, Edit, Selection, View, Go, Run, ...
- Title Bar:** Project\_OS
- Sidebar (PROJECT OS):** EXPLORER, PROJECT\_OS, app.py, OS\_Documentation\_Proposal.pdf, readers\_writers.ipynb, simulation\_data.csv.
- Current Notebook:** readers\_writers.ipynb
- Code Cell [7]:** Imports threading, time, random, datetime, pandas, IPython.display, ipywidgets. Prints "All libraries imported successfully!" followed by a separator line.
- Output Cell [7]:** Shows the printed message: "All libraries imported successfully!"
- Code Cell [8]:** Defines a class ReadersWritersProblem with \_\_init\_\_ method. It initializes reader\_count to 0, waiting\_writers to 0, and two locks: resource\_lock and mutex.
- Output Cell [8]:** Shows the definition of the ReadersWritersProblem class.
- Bottom Navigation:** PROBLEMS, OUTPUT, DEBUG CONSOLE, TERMINAL, PORTS, JUPYTER, python, Cell 4 of 6, 12:03 PM, 1/8/2026.

## Appendix B: Installation Guide

### System Requirements:

- ✓ Python 3.8 or higher
- ✓ 4GB RAM minimum
- ✓ Modern web browser (Chrome, Firefox, Edge)

### Installation Steps:

```
# 1. Install Python packages: pip install streamlit pandas jupyter ipywidgets  
# 2. Run Streamlit appstreamlit run app.py  
# 3. Run Jupyter notebookjupyter notebook readers_writers.ipynb
```

## Appendix C: Troubleshooting

### Issue 1: Jupyter command not found

```
python -m jupyter notebook
```

### Issue 2: Streamlit command not found --- use this command to run project

```
python -m streamlit run app.py Issue
```

### 3: Port already in use streamlit run

```
app.py --server.port 8502
```

## Appendix D: Project Structure

```
OS_Project/
    └── app.py                  # Streamlit application
    └── readers_writers.ipynb   # Jupyter notebook
    └── requirements.txt        # Python dependencies
    └── README.md               # Project overview
    └── documentation.pdf      # This document
```

\* \* \* \* \*